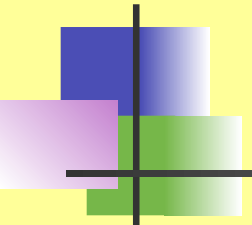




Uninformed Search



Measuring problem-solving performance

Completeness: Is the algorithm guaranteed to find a solution when there is one?

Optimality: Does the strategy find the optimal solution, as defined on page 68?

Time complexity: How long does it take to find a solution?

Space complexity: How much memory is needed to perform the search?

Uninformed search strategies

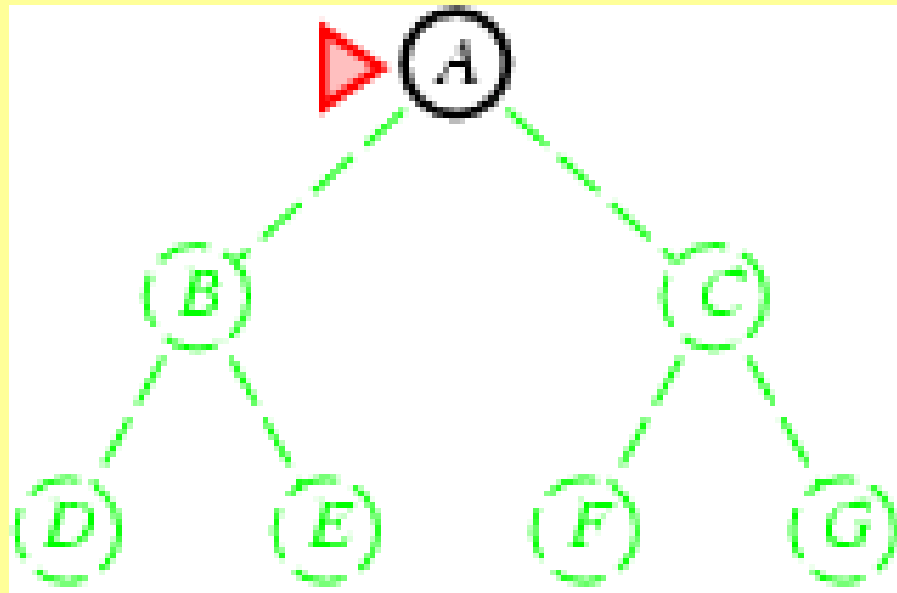


- **Uninformed:** While searching you have no clue whether one non-goal state is better than any other. Your search is blind.
- **Various blind strategies:**
 - Breadth-first search
 - Uniform-cost search
 - Depth-first search
 - Iterative deepening search

Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a first-in-first-out (FIFO) queue, i.e., new successors go at end of the queue.

Is A a goal state?




Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end

Expand:
fringe = [B,C]

Is B a goal state?



Breadth-first search


Expand shallowest unexpanded node

- Implementation:

- *fringe* is a FIFO queue, i.e., new successors go at end

Expand:
fringe=[C,D,E]

Is C a goal state?





Breadth-first search

- Expand shallowest unexpanded node
- Implementation:
 - *fringe* is a FIFO queue, i.e., new successors go at end

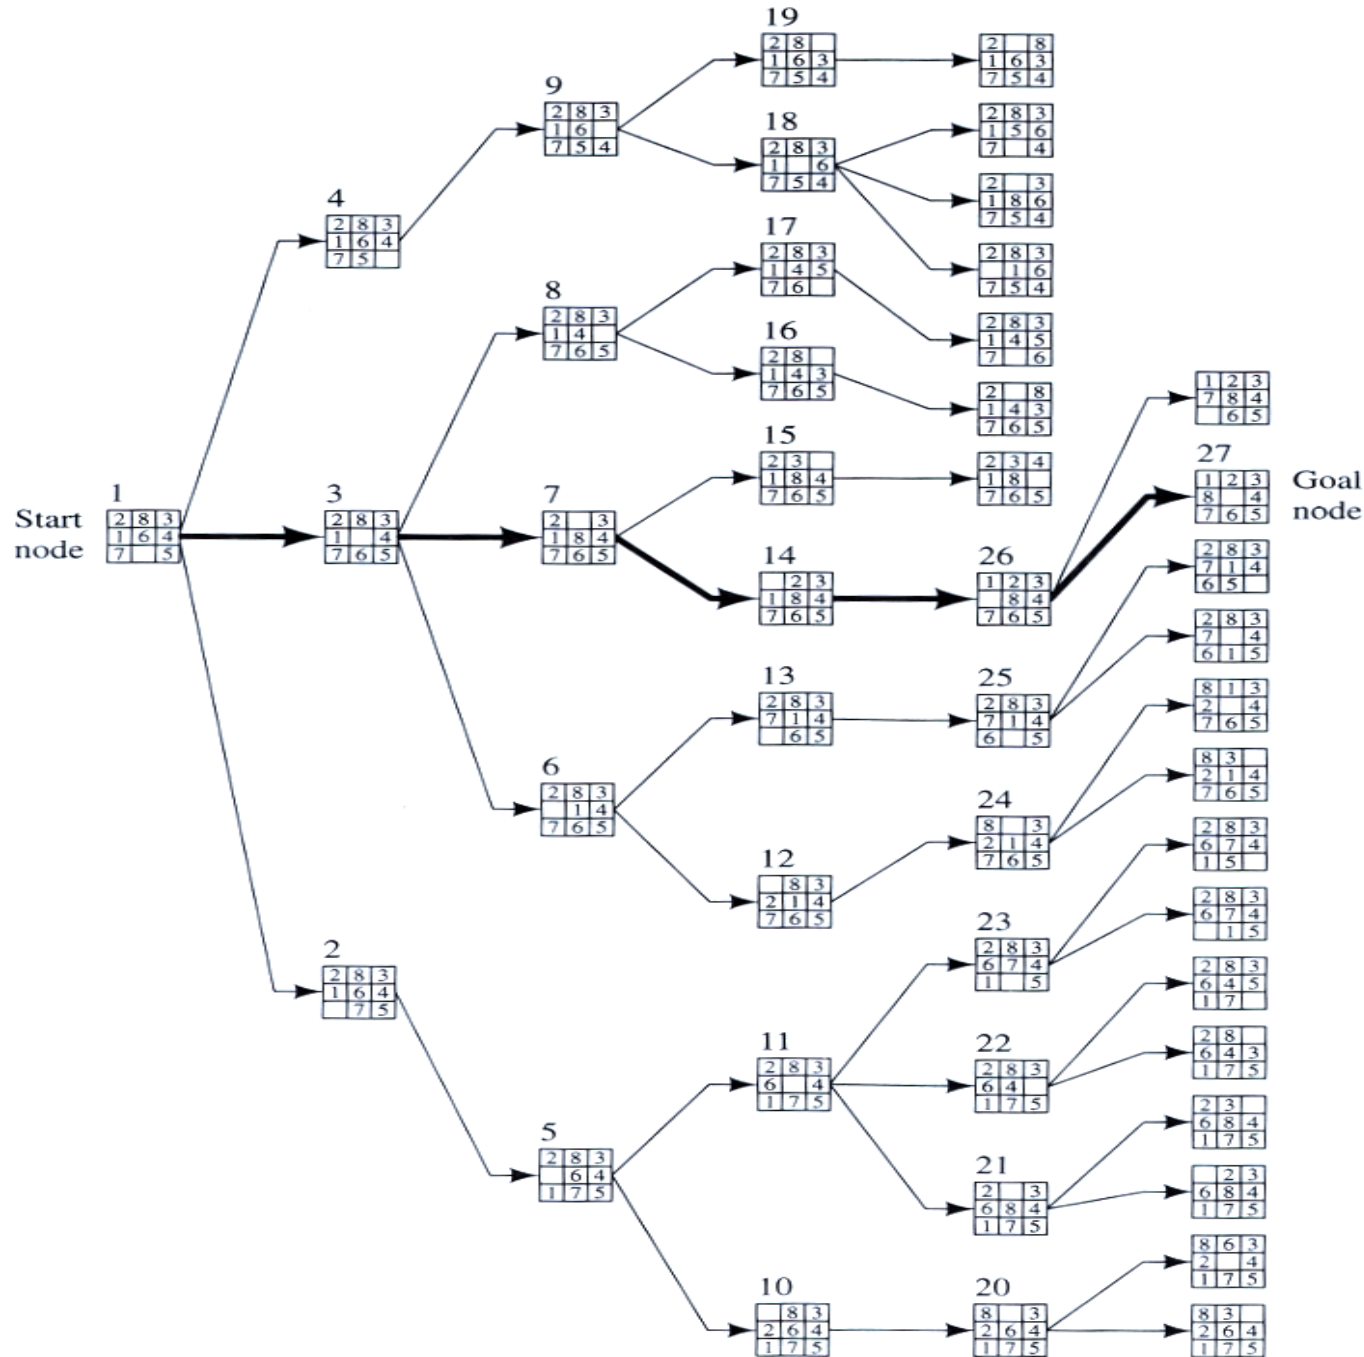
Expand:
fringe=[D,E,F,G]

Is D a goal state?





Example BFS



Properties of breadth-first search

- Complete? Yes it always reaches goal (if b is finite)
- Time? $1+b+b^2+b^3+\dots +b^d + (b^{d+1}-b)) = O(b^{d+1})$
(this is the number of nodes we generate)
- Space? $O(b^{d+1})$ (keeps every node in memory, either in fringe or on a path to fringe).
- Optimal? Yes (if we guarantee that deeper solutions are less optimal, e.g. step-cost=1).
- **Space** is the bigger problem (more than time)

Uniform-cost search

Breadth-first is only optimal if step costs is increasing with depth (e.g. constant). Can we guarantee optimality for any step cost?

Uniform-cost Search: Expand node with smallest path cost $g(n)$.




Figure 3.13 A route-finding problem. (a) The state space, showing the cost for each operator. (b) Progression of the search. Each node is labelled with $g(n)$. At the next step, the goal node with $g = 10$ will be selected.



Uniform-cost search

Implementation: *fringe* = queue ordered by path cost
Equivalent to breadth-first if all step costs all equal.

Complete? Yes, if step cost $\geq \epsilon$
(otherwise it can get stuck in infinite loops)

Time? # of nodes with *path cost* \leq cost of optimal solution.


Space? # of nodes on paths with path cost \leq cost of optimal solution.

Optimal? Yes, for any step cost.

Depth-first search

- Expand *deepest* unexpanded node
- Implementation:
 - *fringe* = Last In First Out (LIPO) queue, i.e., put successors at front

Is A a goal state?



Depth-first search


Expand deepest unexpanded node

- Implementation:

- $fringe = L$

queue=[B,C]

Is B a goal state?




Depth-first search

Expand deepest unexpanded node

- Implementation:

- *fringe* = LIFO queue, i.e., put successors at front
queue=[D,E,C]

Is D = goal state?



Depth-first search


Expand deepest unexpanded node

- **Implementation:**

- *fringe* = LIFO queue, i.e., put successors at front

queue=[H,I,E,C]

Is H = goal state?




Depth-first search

Expand deepest unexpanded node

- Implementation:

- *fringe* = LIFO queue, i.e., put successors at front
queue=[I,E,C]

Is I = goal state?



Depth-first search


Expand deepest unexpanded node

- Implementation:

- $fringe = L$

queue=[E,C]

Is E = goal state?



Depth-first search


Expand deepest unexpanded node

- Implementation:

- $fringe = L$

queue=[J,K,C]

Is J = goal state?



Depth-first search


Expand deepest unexpanded node

- Implementation:

- $fringe = L$

queue=[K,C]

Is K = goal state?



Depth-first search


Expand deepest unexpanded node

- Implementation:

- $fringe = L$

queue=[C]

Is C = goal state?



Depth-first search


Expand deepest unexpanded node

- Implementation:

- $fringe = L$

queue=[F,G]

Is F = goal state?



Depth-first search


Expand deepest unexpanded node

- Implementation:

- $fringe = L$

queue=[L,M,G]

Is L = goal state?



Depth-first search


Expand deepest unexpanded node

- Implementation:


- $fringe = L$

queue=[M,G]


Is M = goal state?




Example DFS



(a)



(b)



(c)



Properties of depth-first search

- Complete? No: fails in infinite-depth spaces
Can modify to avoid repeated states along path
- Time? $O(b^m)$ with m =maximum depth
- terrible if m is much larger than d
 - but if solutions are dense, may be much faster than
breadth-first
- Space? $O(bm)$, i.e., linear space! (we only need to remember a single path + expanded unexplored nodes)
- Optimal? No (It may find a non-optimal goal first)



Iterative deepening search

- To avoid the infinite depth problem of DFS, we can decide to only search until depth L , i.e. we don't expand beyond depth L .
 - ☾ **Depth-Limited Search**
- What of solution is deeper than L ? ☾ Increase L iteratively.
 - ☾ **Iterative Deepening Search**
- As we shall see: this inherits the memory advantage of Depth-First search.

Iterative deepening search $L=0$

Limit = 0



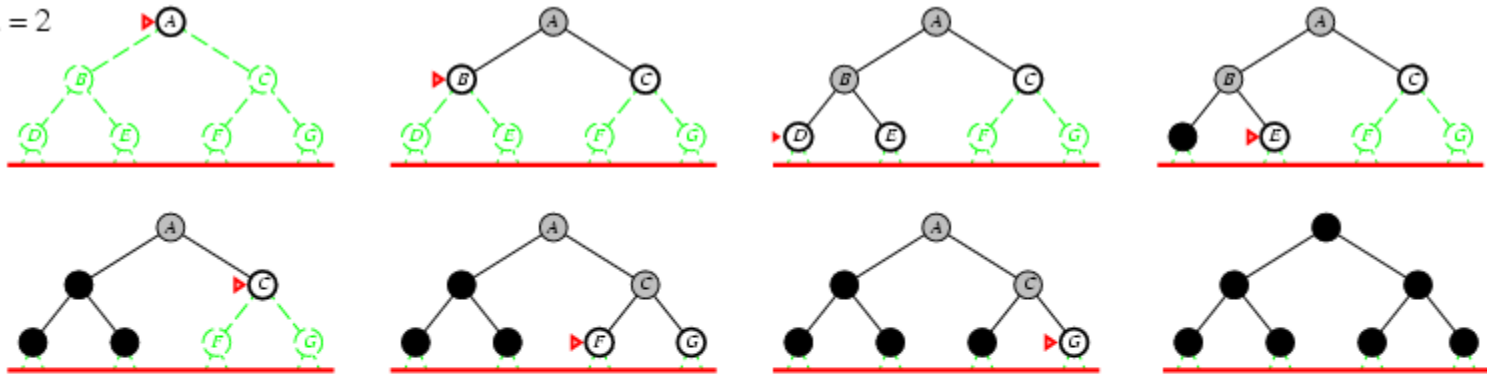
Iterative deepening search $L=1$

Limit = 1




Iterative deepening search $L=2$

Limit = 2



Iterative deepening search $L=3$

Limit = 3



Iterative deepening search

- Number of nodes generated in a depth-limited search to depth d with branching factor b :

$$N_{DLS} = b^0 + b^1 + b^2 + \dots + b^{d-2} + b^{d-1} + b^d$$

- Number of nodes generated in an iterative deepening search to depth d with branching factor b :

$$N_{IDS} = (d+1)b^0 + d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d =$$
$$O(b^d) \neq O(b^{d+1})$$

- For $b = 10, d = 5$,

- $N_{DLS} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$
- $N_{IDS} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$
- $N_{BFS} = \dots = 1,111,100$


BFS




Properties of iterative deepening search

- Complete? Yes
- Time? $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- Space? $O(bd)$
- Optimal? Yes, if step cost = 1 or increasing function of depth.


Example IDS




Depth bound = 1



Depth bound = 2



Depth bound = 3



Depth bound = 4

Stages in Iterative-Deepening Search



Bidirectional Search

- Idea
 - simultaneously search forward from S and backwards from G
 - stop when both “meet in the middle”
 - need to keep track of the intersection of 2 open sets of nodes
- What does searching backwards from G mean
 - need a way to specify the predecessors of G
 - this can be difficult,
 - e.g., predecessors of checkmate in chess?
 - what if there are multiple goal states?
 - what if there is only a goal test, no explicit list?

Bi-Directional Search

Complexity: time and space complexity are: $O(b^{d/2})$




Fig. 2.10 Bidirectional and unidirectional breadth-first searches.




Summary of algorithms


Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

Repeated states


- Failure to detect repeated states can turn a linear problem into an exponential one!



Solutions to Repeated States



State Space



Example of a Search Tree

- Method 1 ← suboptimal but practical
 - do not create paths containing cycles (loops)
- Method 2 ← optimal but memory inefficient
 - never generate a state generated before
 - must keep track of all possible states (uses a lot of memory)
 - e.g., 8-puzzle problem, we have $9! = 362,880$ states



Summary

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed search strategies
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms

<http://www.cs.rmit.edu.au/AI-Search/Product/>

<http://aima.cs.berkeley.edu/demos.html> (for more demos)