


ECE 595
Computer Network Systems


TCP
Reliable Transmission

Transport Layer



Transport Protocols Concern only End Hosts, not Routers


- Header generated by sender is interpreted only by the destination
- Routers view transport header as part of the payload



Transport Layer in Internet

- Purpose 1: (De)multiplexing of data streams to different application processes
- Purpose 2: Provide value-added services that many applications want
 - Recall network layer in Internet provides a “Best-effort” service only, transport layer can add value to that
 - Application may want reliability, etc
 - No need to reinvent the wheel each time you write a new application

Using Transport Layer Port Number to (De)multiplex traffic




In TCP, a data stream is identified by a set of numbers:
(Source Address, Destination Address, Source Port, Destination Port)

Popular Transport Protocols

- UDP:
 - Barebones,minimal
 - Does not provide much functionality besides multiplexing
- TCP:
 - Elaborate, lots of additional functionality provided.

User Datagram Protocol (UDP)

- Connectionless datagram
 - Socket: SOCK_DGRAM
- Port number used for (de)multiplexing
 - port numbers = connection/application endpoint
- Adds end-to-end reliability through optional checksum
 - protects against data corruption errors between source and destination (links, switches/routers, bus)
 - does not protect against packet loss, duplication or reordering



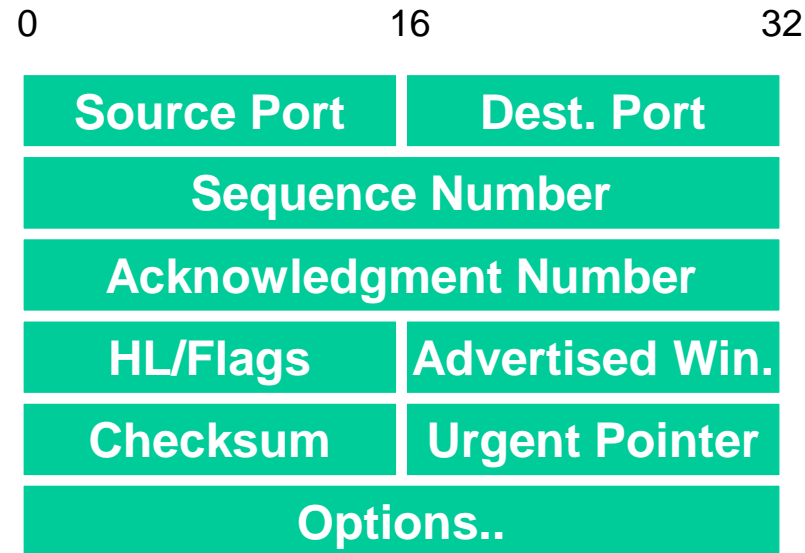
Using UDP

- Custom protocols/applications can be implemented on top of UDP
 - use the port addressing provided by UDP
 - implement own reliability, flow control, ordering, congestion control as it sees fit
- Examples:
 - remote procedure call
 - Multimedia streaming (real time protocol)
 - distributed computing communication libraries

Transmission Control Protocol (TCP)

- Reliable bidirectional in-order byte stream
 - Socket: SOCK_STREAM
- Lots of functionality
- Connection establishment.
 - Logical end-to-end connection, connection state to optimize performance
- Error control
 - Hide unreliability of the network layer from applications
 - Many types of errors: corruption, loss, duplication, reordering.
- End-to-end flow control and congestion control
 - Avoid flooding the receiver and network.

TCP Header




- 20 bytes total
- Sequence Number, ACK, Advertised window: relate to TCP functionality for achieving reliable delivery.
- Sequence Number is for forward direction, ACK for reverse direction.
- HL: Specifies Header Length
- Flags: 6 flags in all
- Urgent pointer: not common TCP usage, used to signal certain data is “out-of-band” and must be processed immediately

Important TCP Flags

- SYN: Synchronize
 - Used when setting up connection
- FIN: Finish
 - Used when tearing down connection
- ACK
 - Acknowledging received data
- RESET:
 - Receiver wants to abort connection, as it received unexpected segment.
- Push and Urgent flags:
 - Not as commonly used.
 - Signify receiving process must be notified, or out-of-band data

Bidirectional Communication



- Each Side of Connection can Send *and* Receive
- What this Means
 - Maintain different sequence numbers for each direction
 - Single segment can contain new data for one direction, plus acknowledgement for other
 - But some contain only data & others only acknowledgement


Ongoing Communication

- Bidirectional Communication
 - Each side acts as sender & receiver
 - Every message contains acknowledgement of received sequence
 - Even if no new data have been received
 - Every message advertises window size
 - Size of its receiving window
 - Every message contains sent sequence number
 - Even if no new data being sent
- When Does Sender Actually Send Message?
 - When sending buffer contains at least max. segment size (- header sizes) bytes
 - When application tells it
 - Set PUSH flag for last segment sent
 - When timer expires

Connection Setup

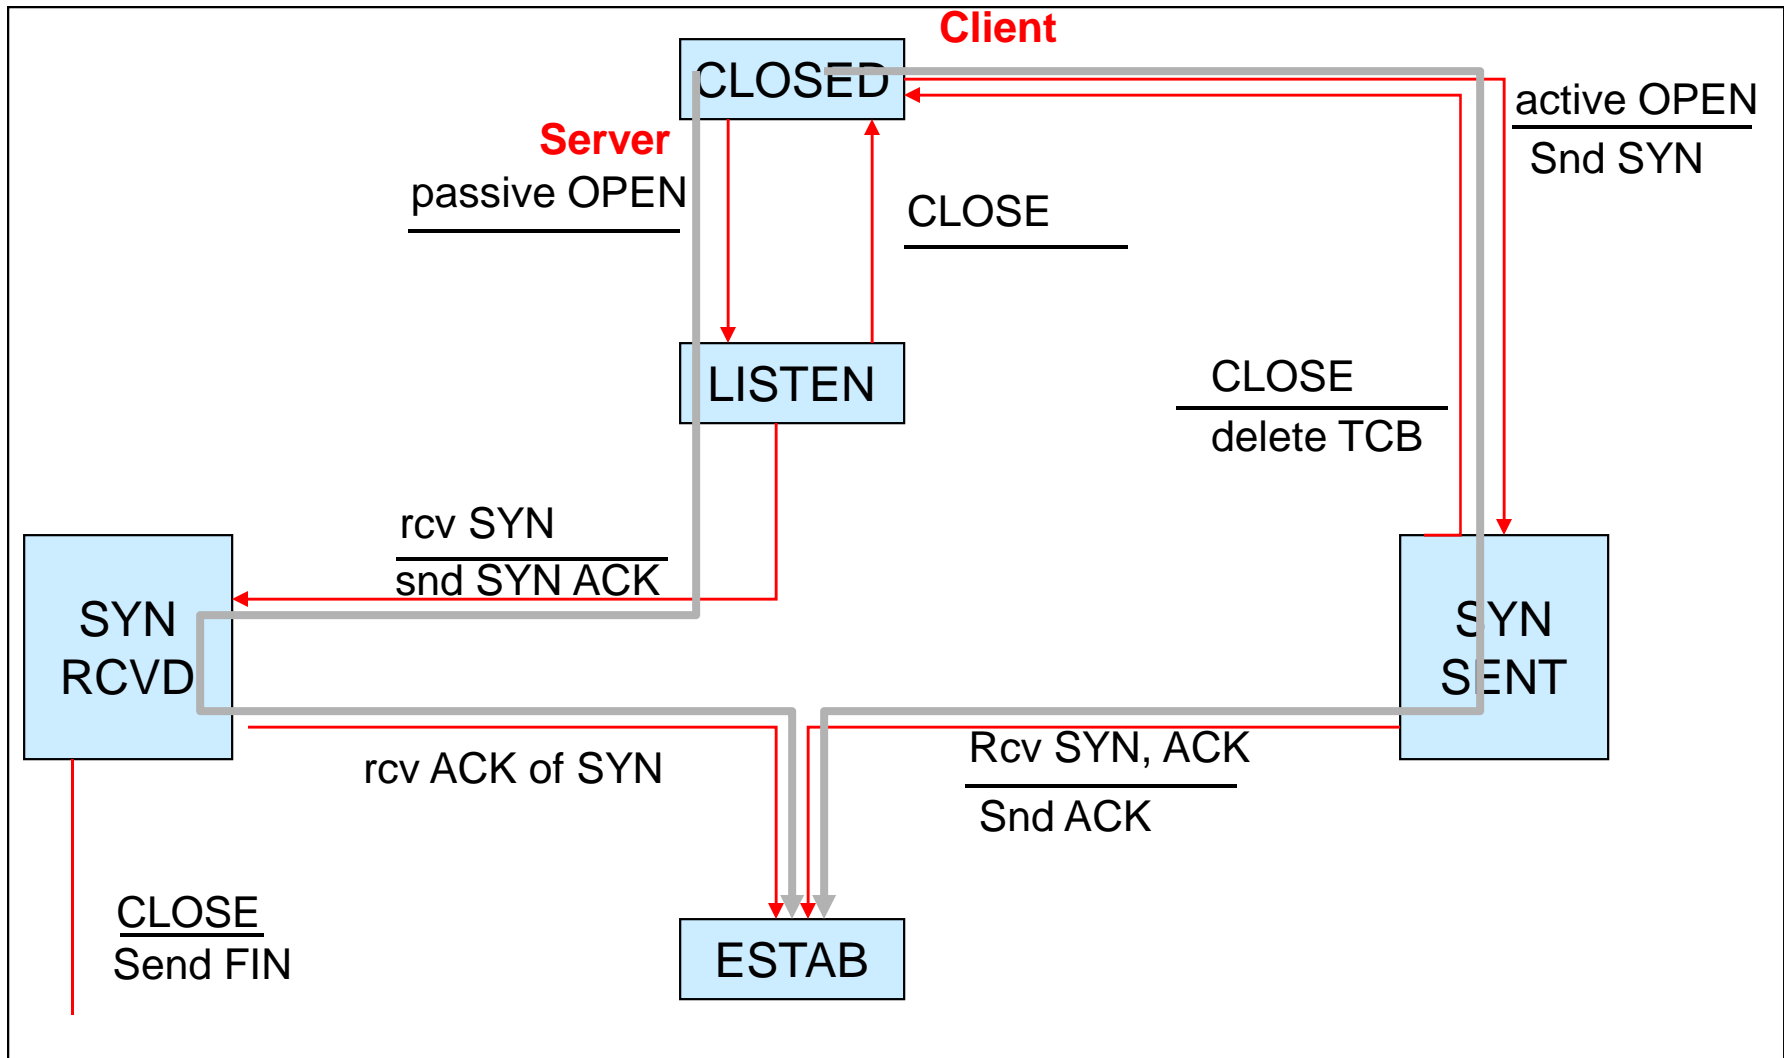
- Why need connection setup?
- Mainly to agree on starting sequence numbers
 - Starting sequence number is randomly chosen
 - Reason, to reduce the chance that sequence numbers of old and new connections from overlapping

Establishing Connection



- Three-Way Handshake
 - Each side notifies other of starting sequence number it will use for sending
 - Each side acknowledges other's sequence number
 - SYN-ACK: Acknowledge sequence number + 1
 - Can combine second SYN with first ACK

TCP State Diagram: Connection Setup

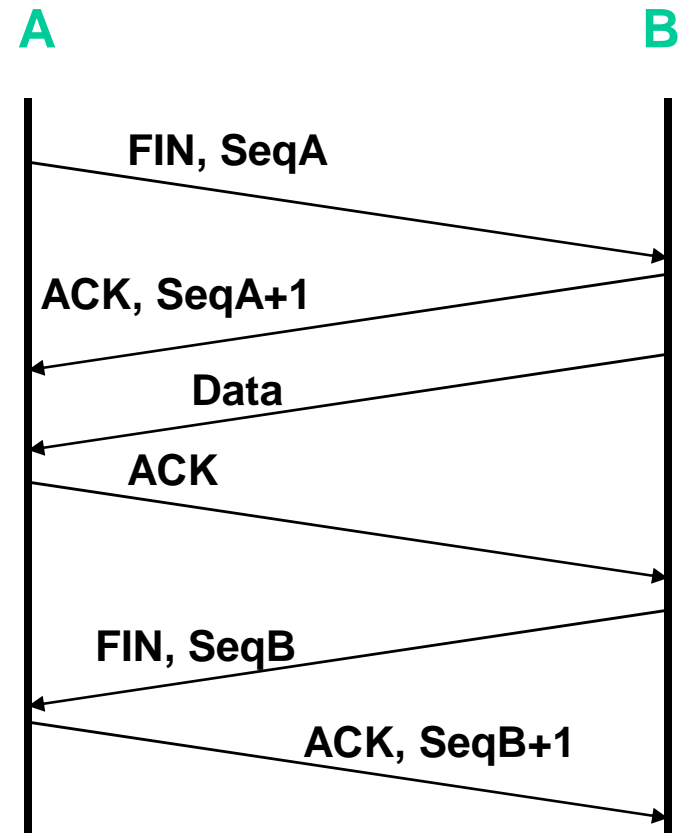


SYN-Flood Attacks

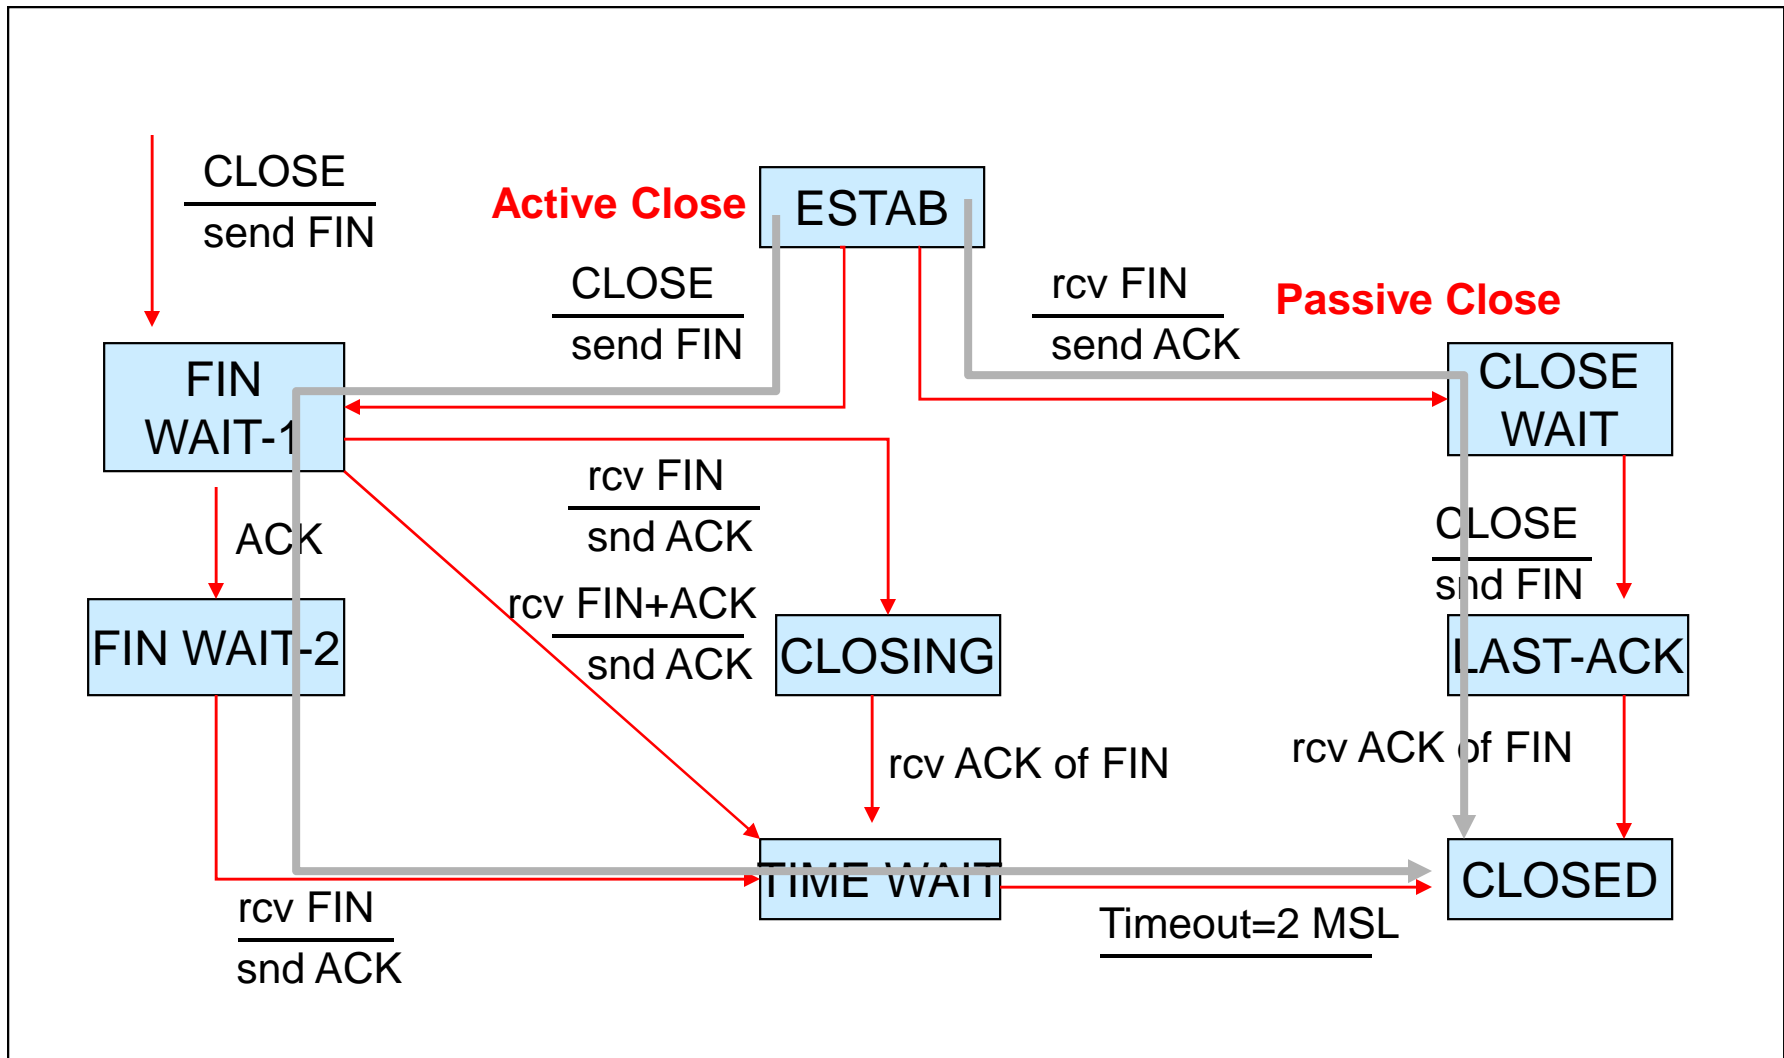
- Exploit 3-way handshake
- Attacker spoofs IP addresses, sends a large number of SYNs to the destination.
- Creates “partial unopened” state at the server
- Server sends SYN-ACK to (spoofed) addresses, to which no response is received.
- Rate at which SYNs sent much faster than rate at which server eliminates partial state.

Tearing Down Connection

- Either Side Can Initiate Tear Down
 - Send FIN signal
 - “I’m not going to send any more data”
- Other Side Can Continue Sending Data
 - Half open connection
 - Must continue to acknowledge
- Acknowledging FIN
 - Acknowledge last sequence number + 1



State Diagram: Connection Tear-down




Reliable Transmission

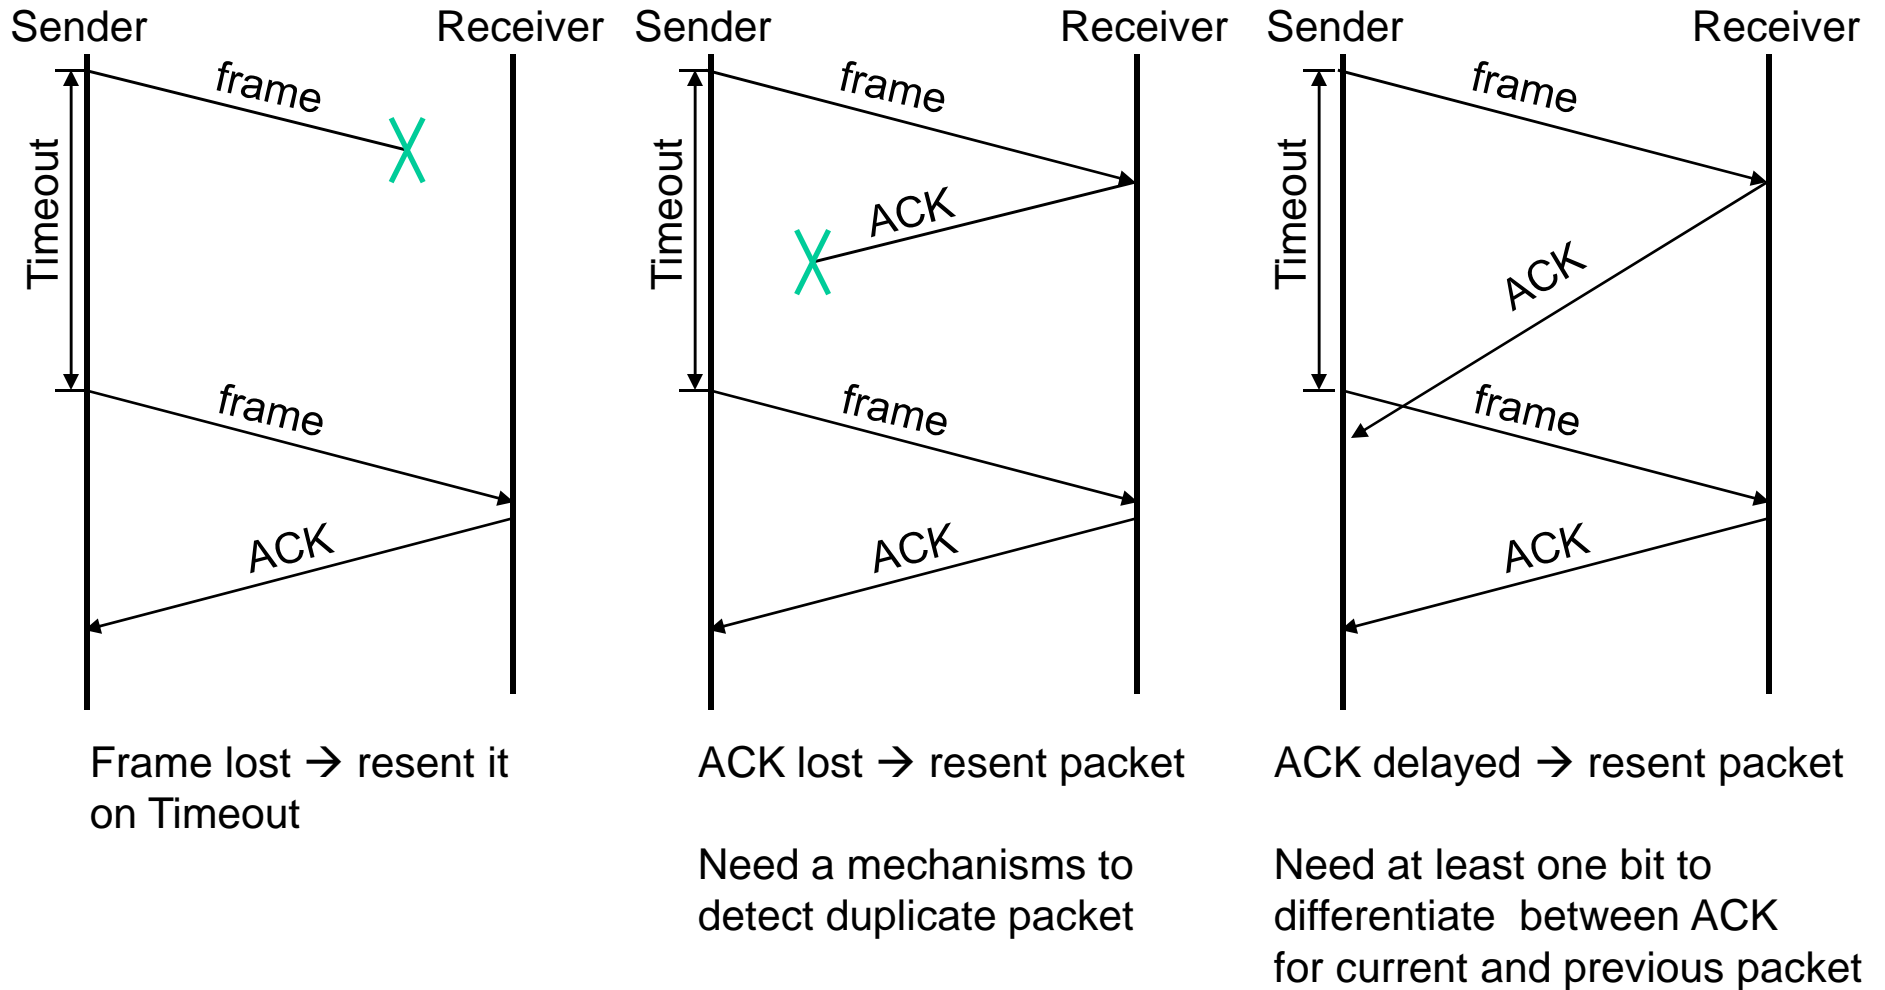
- Use two basic techniques:
 - Acknowledgements (ACKs)
 - Timeouts
- Two examples:
 - Stop-and-go
 - Sliding window

Stop-and-Go

- Receiver: send an acknowledge (ACK) back to the sender upon receiving a packet (frame)
- Sender: excepting first packet, send a packet only upon receiving the ACK for the previous packet




What Can Go Wrong?




Stop-and-Go Disadvantage

- May lead to inefficient link utilization
- Example: assume
 - One-way propagation = 15 ms
 - Bandwidth = 100 Mbps
 - Packet size = 1000 bytes \rightarrow transmit = $(8 \cdot 1000) / 10^8 = 0.08 \text{ ms}$
 - Neglect queue delay \rightarrow Latency = approx. 15 ms; RTT = 30 ms




Stop-and-Go Disadvantage (cont'd)

- Send a message every 30 ms \rightarrow Throughput = $(8 \times 1000) / 0.03 = 0.2666$ Mbps
- Thus, the protocol uses less than 0.3% of the link capacity!



Sliding Window Approach




Sliding window protocol

Send multiple packets without waiting for ACK.

Sliding window size: Max number of packets that can be sent without ACK being received (3 in figure)


Sequence Numbers in TCP

- Each byte in byte stream is numbered.
 - 32 bit value
- TCP breaks up the byte stream in segments
- Each segment has a sequence number.
 - Indicates where it fits in the byte stream




Sliding Window Protocol: Sender

- TCP: operates at a byte stream level (rather than packet level)
- Sender maintains a window of sequence numbers
 - SWS (sender window size) – maximum data that can be sent without receiving an ACK
 - LBA (last byte acknowledged)
 - LBS (last byte sent)
- TCP sender side socket buffer:
 - Data sent but not acknowledged
 - Data written by application but not sent yet




Sliding Window Protocol: Receiver

- TCP maintains a receive socket buffer
 - Stores data that arrived out-of-order (cannot be given to application yet)
 - Also stores data that arrived in-order but not yet read by application (slow process)
- Receiver maintains a window of sequence numbers
 - NBE (next byte expected – all previous bytes received in sequence)
 - LBR (last byte read by application)
 - LBG (last byte got by receiver)



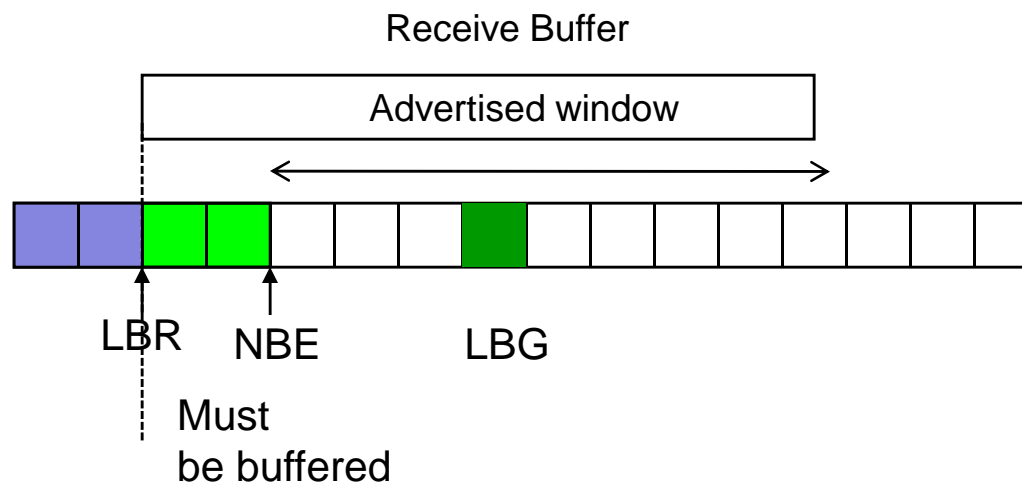
Sliding Window Protocol: Receiver

- If incoming byte $< \text{NBE}$
 - Discard packet
- Else
 - Accept packet (provided it fits in the buffer)
 - ACK largest byte that all previous bytes were received



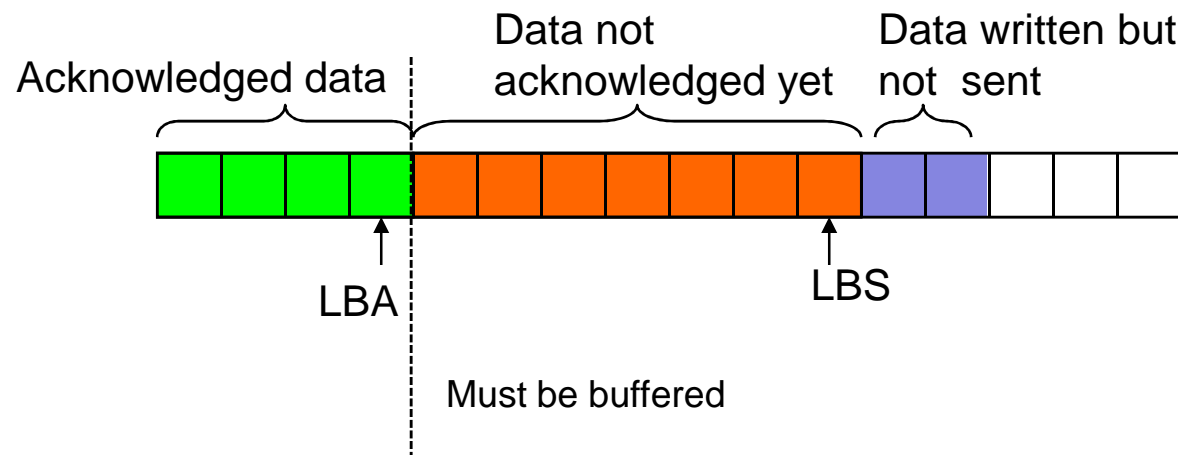
Flow Control

- Receiver throttles the sender by advertising a window no larger than the amount of data it can buffer
- Advertised Window depends on
 - Size of receive buffer
 - How fast receiver process can read the bytes from the buffer
- TCP sending window impacted by advertised window



Flow Control (sender side)

- Receiver throttles the sender by advertising a window no larger than the amount of data it can buffer
- Effective window at sender (limits how much it can send)
Advertised window – (outstanding unacknowledged data)
- Sending process can write to send-side TCP buffer only if space in sender side buffer – otherwise will block



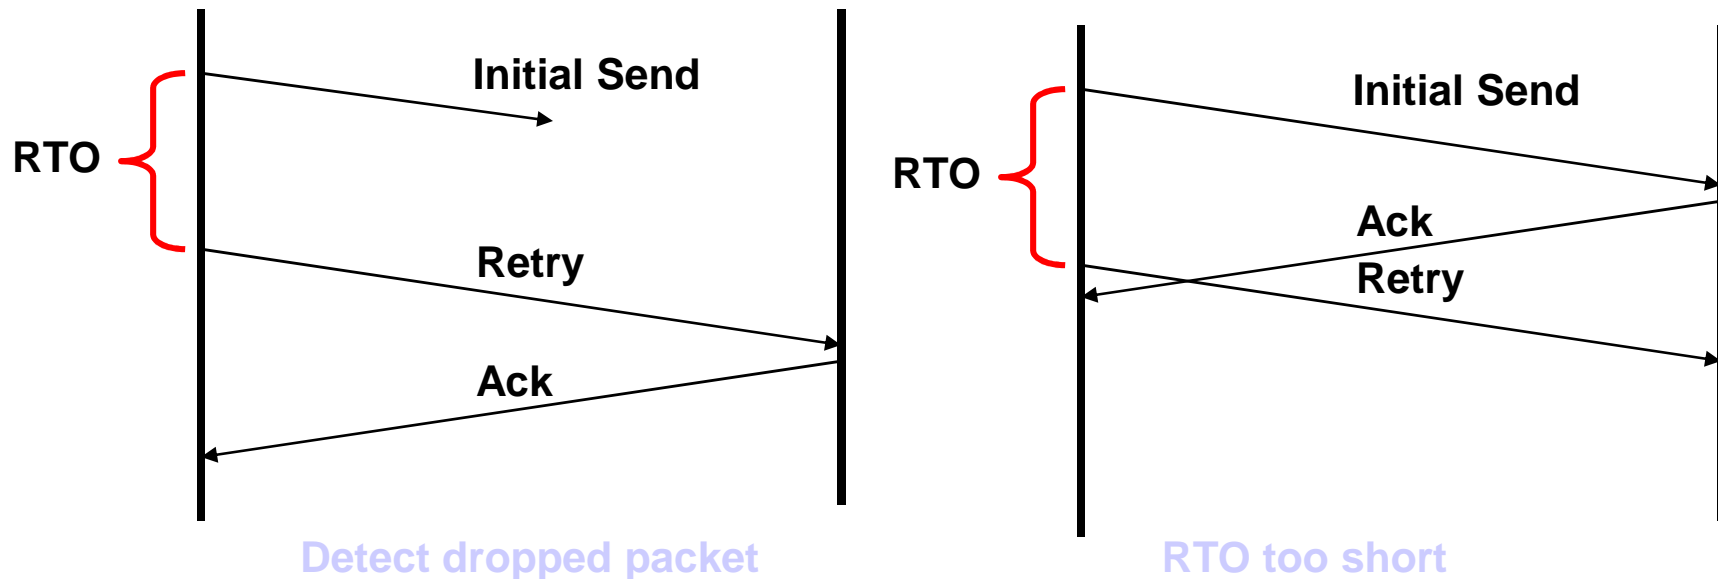
Buffer sizes

- Size of TCP socket buffers at sender and receiver side critical
- Would like this to be Bandwidth * Delay product of network, to ensure the pipe can be full.
- Must be tuned carefully for individual path.
- Traditionally, OSs used default socket buffer sizes that must be overridden by application programmer/end-user
 - “Tuning” in manual fashion tricky/hard.
- Research efforts at “auto-tuning” TCP buffers:
 - [Automatic TCP Buffer Tuning](#)
Jeffrey Semke, Matthew Mathis, and Jamshid Mahdavi, SIGCOMM 1998
- More recently finding its way into some Operating Systems.

Sequence Number Wraparound

- TCP sequence numbers: 32 bits.
- Could wraparound
- Assuming packets remain in the network at most 120s, wrap around is a problem only for bandwidth above 286 Mbps
 - OK for most common networks
 - An issue for high-speed networks (e.g. special purpose scientific networks)
 - 120s: “Maximum segment lifetime”: just a conservative estimate of how long a packet may linger in the network: not actual value.

Setting Retransmission Timeout (RTO)




- Time between sending & resending segment
- Challenge
 - Too long: Add latency to communication when packets dropped
 - Too short: Send too many duplicate packets
 - General principle: Must be > 1 Round Trip Time (RTT)

Setting Timeouts

- Retransmission time-out should be set based on round-trip delay
- But round-trip delay different for each path!
- Must estimate RTT dynamically

Round-trip Time Estimation


- Every Data/Ack pair gives new RTT estimate



- Can Get Lots of Short-Term Fluctuations

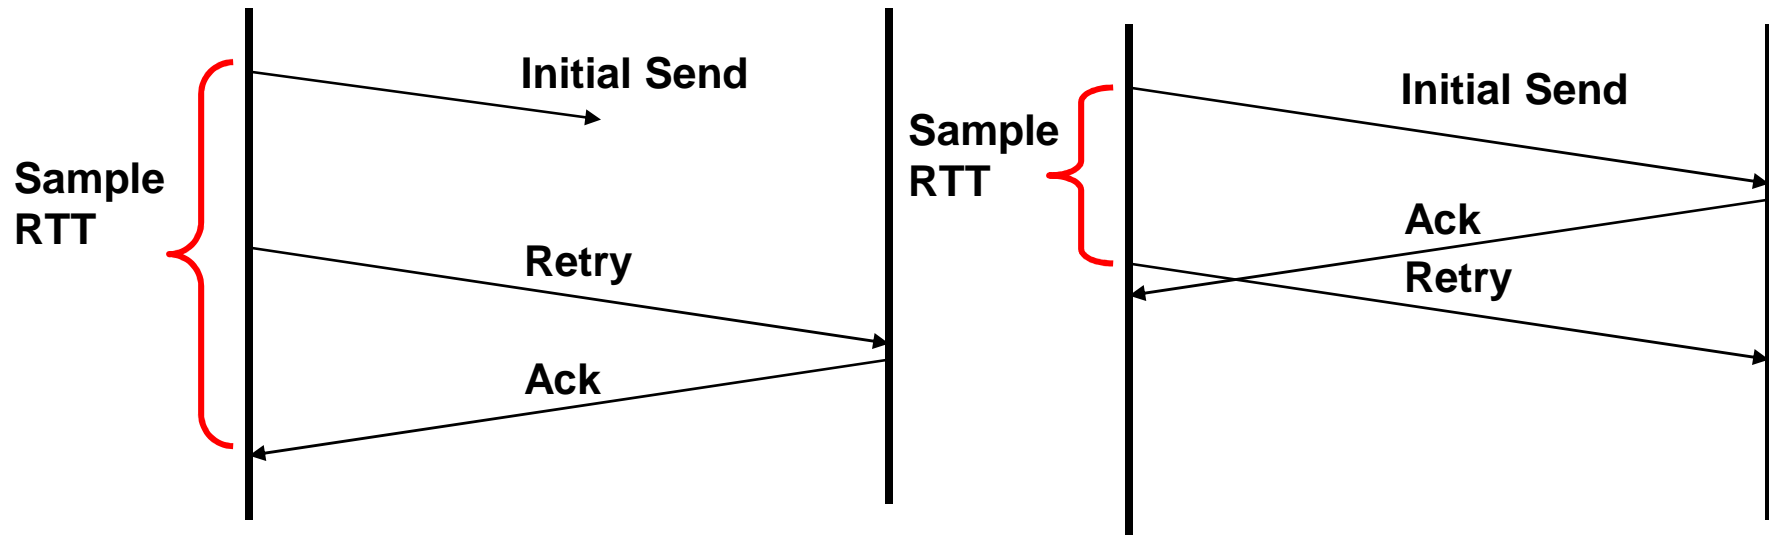
Original TCP Round-trip Estimator

- Round trip times exponentially averaged:
 - New RTT = α (old RTT) + $(1 - \alpha)$ (new sample)
 - Recommended value for α : 0.8 - 0.9
 - 0.875 for most TCP's



- Retransmit timer set to $(b * RTT)$, where $b = 2$
 - Every time timer expires, RTO exponentially backed-off
- Enhanced algorithm used in practice
 - [Jacobson/Karels: see textbook if interested]

Issue



What value of Sample RTT to use?

Solution: TCP simply neglects SampleRTT in such cases.
[Karn/Partridge algorithm]